

Process Management

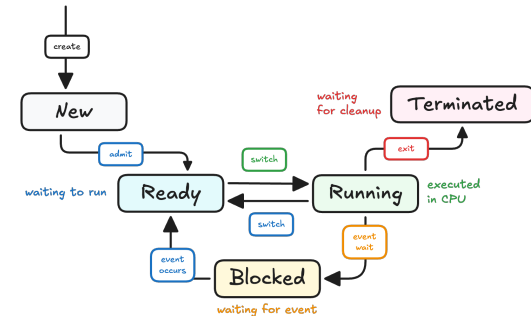
Stack frame structure:

Local variables, parameters, return PC, other info

Setup/Teardown stack frame:

1. **Caller** passes parameters with register and/or stack
2. **Caller** saves return PC on stack
3. **Callee** saves registers used, old FP, SP
4. **Callee** allocates space (callee local variables) stack.
5. **Callee** adjusts SP to point to new stack top.
6. Execution of call
7. **Callee** restores saved registers, FP, SP
8. **Callee** places return result on stack
9. **Callee** restores saved SP
10. **Caller** utilises return result
11. **Caller** continues execution

5-State Process Model



Syscall mechanism

1. User program: **library call**
2. **Library call**: syscall no. of designated loc.
3. **Library call**: TRAP instruction to kernel mode.
4. **System call** handler is determined using syscall.
5. **Syscall** handler is executed.
6. Control return to **lib call**, switch back to user mode.
7. **Library call** returns to user program.

Process Control Block: information about a process (registers, memory region information, PID, process state)

Exception vs Interrupt:

Synchronous vs **asynchronous**, executes **exception handler** vs executes **interrupt handler**

Abstractions in UNIX

pid_t fork(void): Returns: Parent: **child PID**, child: **0**.
 int exec *(...) : Replace code with another
 int wait(int *status): Waits for process to end

Process Scheduling

Non-preemptive: Process stays scheduled

Preemptive: Process is suspended at time quota

Batch processing

Turnaround time Total time from arrival to finish

Throughput Number of tasks/unit time

CPU utilisation % of time CPU is doing work

Scheduling algorithms

FCFS: FIFO queue, **No starvation**

SJF: Smallest CPU time, **Possible starvation**, predicts using exponential average of history.

SRT: Preemptive SJF

Convoy effect Fight for CPU and fight for I/O together

Interactive environment

Response time: Time between request & response

Scheduling algorithms

RR: FCFS, preemptive

Priority: Each task gets priority, highest first

preemptive: new process preempts, **non-preemptive**: wait for next round. can result in starvation

Priority inversion: higher priority task forced to block

MLFQ:

$p(A) > p(B)$, A runs — $p(A) == p(B)$, RR.

New job gets highest priority. If it uses time slice fully, priority **reduced**. Else, **retained**.

Lottery: Tickets assigned, randomly chosen winner

Threads

User thread: Implemented as user library

More flexible, **kernel unaware**

Kernel thread: Implemented in OS

Thread level scheduling possible, **slower**, **less flexible**

Hybrid thread: Both

```

int pthread_create(pthread_t *thread,
  const pthread_attr_t *attr,
  void *(*start_routine) (void *), void *arg);
  
```

```

int pthread_exit(void *retval);
int pthread_join(pthread_t thread, void **retval)
  
```

Inter-Process Communication

Shared memory

Efficient, **easy to use**, **synchronisation**, **harder to implement**

```

int shmget(key_t key, size_t size, int shmflg)
void *shmat(int shmid, const void *shmaddr,
  int shmflg, int shmdt)
int shmdt(const void *shmaddr)
int shmctl(int shmid, int cmd, struct shmid_ds *buf)
  
```

Message parsing

Portable, **easier sync**, **inefficient**, **harder to use**

Direct One buffer per pair of (sender, receiver)

Indirect Send/receive to port/mailbox

Blocking Until message received/arrives

Non-blocking (Async) Does not block

Unix Pipes Fixed size circular byte buffer

Writer wait when buffer full, reader wait when empty

```

int pipe(int pipefd[2])
pipefd[0] \ read, pipefd[1] \ write
  
```

Synchronisation

Properties

Mutual exclusion If process in CS, all other processes prevented from entering

Progress No process in CS, one waiting process should be granted access

Bounded wait Process requests to enter, there should be an upper bound on amount of processes that can enter before it

Independence Process not in CS should not block other process

Challenges

Deadlock All processes blocked

Livelock Processes keep changing state to avoid deadlock

Starvation Processes never get to make progress

Critical Section Implementations

Global Lock = 0

Busy waiting, **Interrupt disabling**, **busy waiting**

Turn-based Want = int[2]

No deadlocks if both Wants are not 1, **deadlock otherwise**

Peterson's Both global and turn-based locks

Assumes turn is atomic.

Busy waiting

Semaphore